

# Simple parallel computation

Use spot instances to run parallel jobs with detailed resource monitoring

Angus Gray-Weale

2024-10-04

*This repository is a work in progress. It has been through limited testing, and requires clearer control of its settings.*

This repository provides a set of **bash** scripts that handles simple parallel execution using GNU Parallel, for example on AWS Spot instances.

The code has two main uses:

- Run simple parallel calculations on a local or cloud workstation, tracking memory use, and load averages.  
Here a simple parallel calculation is the trivial case of independent jobs without communication.
- Run the same type of job on an AWS Spot Instance, checking periodically for interruption notices as advised by the AWS documentation.

Good use is made of **rclone**, GnuPG, GNU Parallel, and GNU Niceload, as explained below.

For use the code must be modified to include code that runs the jobs. Unmodified, the **run.sh** script uses **stress** and **dd** to run a dummy job and generate dummy outputs that aid in design and debugging.

## Possible extensions and improvements

Straightforward extensions of the code include:

- The GNU Parallel command could be replaced by an MPI job running on a single node, allowing the other features of this code to be applied to a parallel job with inter process communication.
- Similar Spot instances are available from other Cloud server farms. With small changes this code could be used with these.
- The code that configures the instance with from a user data script would not be necessary if instead a machine image with all necessary software were used.

## Details of operation

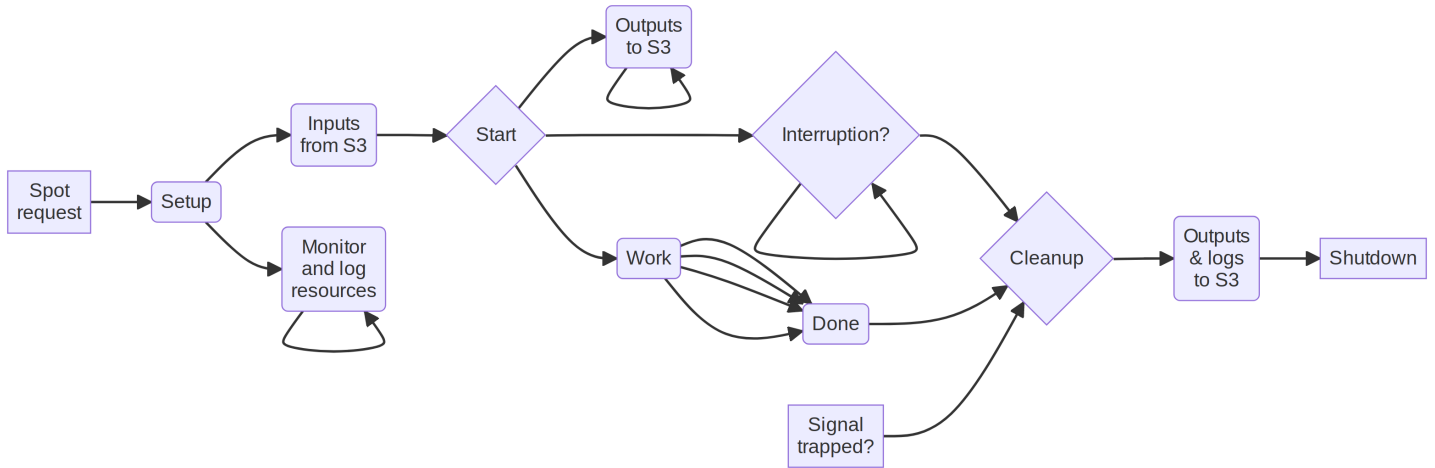


Figure 1: Execution

The above chart shows the flow of execution from left to right. The description here follows execution on a Spot instance. Use on a local machine skips the setup steps and requires only execution of `run.sh`.

1. A **spot request** is issued using a launch template that includes the user data script provided.
2. The user data script installs necessary packages, keys, and this repository. It finishes by executing `run.sh`.
3. A function runs in the background that periodically logs the system's load averages and memory use. The function checks `/dev/shm` for a file containing the ids of all processes started by the computation, and records details of these processes, for example individual memory use.
4. Inputs are then moved from their source to a work space using `rclone`. This might well just copy the files from somewhere else on a local disk, or it might copy them from a remote location such as S3. A configuration file for `rclone` must be provided. The advantage of using `rclone` is that it can run many transfers in parallel and is compatible with a very wide range of data stores.
5. Optionally, and if appropriate keys have been installed, encrypted input files are decrypted.
6. With the system set up, resources monitored, and input files ready, the calculation can begin. Three processes are started:
  - a. A function in the background checks for output files matching a given glob, encrypts them if this has been requested, and then uses `rclone` to copy them to the output destination. This ensures that as many output files as possible have already been saved if an interruption notice is received.
  - b. The main script polls for a spot interruption notice, and if one is found, triggers the cleanup routine.
  - c. GNU Parallel runs in the background, and executes the `run` function for example once for each input file. Extra options may be specified for `parallel` using `OPT_PARALLEL`.
7. The `run` function records processes spawned by the calculation in a file on `/dev/shm`, and applies `niceload` to each so that the system's load is under some control.

8. Signals are trapped and trigger the cleanup routine.
9. Triggering cleanup encrypts the outputs if this has been requested and copies the output files to their destination with `rclone`.
10. If we are using a spot instance the last step is shutdown. The instance should be configured to terminate and delete all disks on shutdown.